



Kafka For Developers Using Spring Boot




Software Development


Past



Current




MicroServices Architecture




What is an Event Streaming Platform?

- Producers and Consumers subscribe to a stream of records




- Store stream of Events



- Analyze and Process Events as they occur

Apache Kafka (Event Streaming Platform)



Traditional Messaging System


- Transient Message Persistence
- Brokers responsibility to keep track of consumed messages
- Target a specific Consumer
- Not a distributed system

Kafka Streaming Platform

- Stores events based on a retention time. Events are Immutable
- Consumers Responsibility to keep track of consumed messages
- Any Consumer can access a message from the broker
- It's a distributed streaming system

Kafka Use Cases


Transportation



Driver-Rider Notifications

Food Delivery Notifications

Retail



Sale Notifications

RealTime Purchase
recommendations

Tracking Online Order
Deliveries


Banking



Fraud Transactions

New Feature/Product
notifications


Kafka Terminology & Client APIs



Kafka Topics & Partitions

Kafka Topics

- Topic is an **Entity** in Kafka with a name




Topic and Partitions

- Partition is where the message lives inside the topic
- Each Topic will be create with one or more partitions


Topic and Partitions

TopicA



- Each Partition is an ordered , immutable sequence of records
- Each record is assigned a sequential number called **offset**
- Each partition is independent of each other
- Ordering is guaranteed only at the partition level
- Partition continuously grows as new records are produced
- All the records are persisted in a commit log in the file system where Kafka is installed



Topics and Partitions



Setting up Kafka in Local




APACHE
ZooKeeper™




Broker registered
with zookeeper

Sending Message Without Key



Sending Message With Key


test-topic



Consumer Groups


- **group.id** is mandatory
- **group.id** plays a major role when it comes to scalable message consumption.

test-topic




group.id = group1 **Consumer 1**

Consumer Groups




Consumer Groups




Consumer Groups

test-topic




Consumer Groups

test-topic




Consumer Groups



Consumer Groups : Summary

- Consumer Groups are used for scalable message consumption
- Each different application will have a unique consumer group
- Who manages the consumer group?
 - Kafka Broker manages the consumer-groups
 - Kafka Broker acts as a Group Co-ordinator

Commit Log




Retention Policy

- Determines how long the message is retained ?
- Configured using the property **log.retention.hours** in **server.properties** file
- Default retention period is **168 hours** (7 days)


What is a Distributed System?

- Distributed systems are a collection of systems working together to deliver a value




Characteristics of Distributed System


- Availability and Fault Tolerance
- Reliable Work Distribution
- Easily Scalable
- Handling Concurrency is fairly easy



Kafka as a Distributed System



Kafka as a Distributed System



- Client requests are distributed between brokers
- Easy to scale by adding more brokers based on the need
- Handles data loss using Replication

Start Kafka Broker

```
./kafka-server-start.sh ../config/server.properties
```

Setting up Kafka Cluster

- New **server.properties** files with the new broker details.


```
broker.id=<unique-broker-id>
listeners=PLAINTEXT://localhost:<unique-port>
log.dirs=/tmp/<unique-kafka-folder>
auto.create.topics.enable=false(optional)
```

Example: **server-1.properties**

```
broker.id=1
listeners=PLAINTEXT://localhost:9093
log.dirs=/tmp/kafka-logs-1
auto.create.topics.enable=false(optional)
```


How Topics are distributed?

```
./kafka-topics.sh -  
-create --topic test-topic-replicated  
-zookeeper localhost:2181  
--replication-factor 3  
--partitions 3
```




How Kafka Distributes Client Requests?

Kafka Producer




How Kafka Distributes Client Requests?

Kafka Consumer



How Kafka Distributes Client Requests?


Kafka Consumer Groups



Summary : How Kafka Distributes the Client Requests?

- Partition leaders are assigned during topic Creation
- Clients will only invoke leader of the partition to produce and consume data
 - Load is evenly distributed between the brokers


How Kafka handles Data loss?




Replication

```
./kafka-topics.sh -  
-create --topic test-topic-replicated  
-zookeeper localhost:2181  
--replication-factor 3  
--partitions 3
```

Replication



Replication



In-Sync Replica(ISR)

- Represents the number of replica in sync with each other in the cluster
 - Includes both **leader** and **follower** replica
- Recommended value is always greater than 1
- Ideal value is **ISR == Replication Factor**
- This can be controlled by **min.insync.replicas** property
 - It can be set at the **broker** or **topic** level

Library Inventory




Library Inventory Flow




Library Inventory




Library Inventory Architecture




Library Inventory Architecture




Library Event Domain



Library Event Domain



Library Events Producer API




KafkaTemplate

Kafka Producer in Spring

KafkaTemplate


- Produce records in to Kafka Topic
 - Similar to JDBCTemplate for DB

How KafkaTemplate Works ?



KafkaTemplate.send()

Behind the Scenes



Configuring KafkaTemplate

Mandatory Values:

bootstrap-servers: localhost:9092,localhost:9093,localhost:9094

key-serializer: org.apache.kafka.common.serialization.IntegerSerializer


value-serializer: org.apache.kafka.common.serialization.StringSerializer

KafkaTemplate AutoConfiguration

application.yml

```
spring:  
  profiles: local  
  kafka:  
    producer:  
      bootstrap-servers: localhost:9092,localhost:9093,localhost:9094  
      key-serializer: org.apache.kafka.common.serialization.IntegerSerializer  
      value-serializer: org.apache.kafka.common.serialization.StringSerializer
```


Library Inventory Architecture



KafkaAdmin

- Create topics Programmatically
- Part of the **SpringKafka**
- How to Create a topic from Code?
 - Create a Bean of type **KafkaAdmin** in SpringConfiguration
 - Create a Bean of type **NewTopic** in SpringConfiguration

Library Events Producer API




PUT - “/v1/libraryevent”

- libraryEventId is a mandatory field

```
{  
  "libraryEventId": 123,  
  "eventStatus": null,  
  "book": {  
    "bookId": 456,  
    "bookName": "Kafka Using Spring Boot",  
    "bookAuthor": "Dilip"  
  }  
}
```


Kafka Producer Configurations

- acks
 - acks = 0, 1 and all
 - acks = 1 -> guarantees message is written to a leader (Default) 
 - acks = all -> guarantees message is written to a leader and to all the replicas 
 - acks=0 -> no guarantee (Not Recommended) 


Kafka Producer Configurations

- retries
 - Integer value = [0 - 2147483647]
 - In Spring Kafka, the default value is -> **2147483647**
- retry.backoff.ms
 - Integer value represented in milliseconds
 - Default value is 100ms

Library Events Consumer



Kafka Consumer



Spring Kafka Consumer

- MessageListenerContainer
 - KafkaMessageListenerContainer
 - ConcurrentMessageListenerContainer
- **@KafkaListener** Annotation
 - Uses ConcurrentMessageListenerContainer behind the scenes

KafkaMessageListenerContainer

- Implementation of MessageListenerContainer
- Polls the records
- Commits the Offsets
- Single Threaded

@KafkaListener


- This is the easiest way to build Kafka Consumer
- KafkaListener Sample Code

```
@KafkaListener(topics = {"${spring.kafka.topic}"})
public void onMessage(ConsumerRecord<Integer, String> consumerRecord) {
    log.info("OnMessage Record : {}", consumerRecord);
}
```


- Configuration Sample Code

```
@Configuration
@EnableKafka
@Slf4j
public class LibraryEventsConsumerConfig {
```

Library Events Consumer




Library Events Producer API




Kafka Producer Errors

- Kafka Cluster is not available
- If **acks= all** , some brokers are not available
- **min.insync.replicas** config
 - Example : **min.insync.replicas = 2**, But only one broker is available




min.insync.replicas



min.insync.replicas = 2

Kafka Cluster



Broker 1




Broker 2



Broker 2

Retain/Recover Failed Records



Retain/Recover Failed Records

Producer Misconfiguration - Option 2

